

# Structurally Valid, Semantically Wrong: Measuring the Unenforced-Rule Gap at a Web Application’s Input Boundary

Zahid Rahimoon

*Independent Researcher*

Email: zahidrahimoon22@gmail.com

**Abstract**—A web application’s schema validation checks that an input is well formed. It does not check the business rules that decide whether the input is appropriate, and in the application studied here most of those rules are not enforced in code. This paper measures how often inputs that pass validation break such rules, and whether inputs written by large language models (LLMs) differ from inputs produced in other ways. A Next.js testbed with an enumerated reference of 43 business rules received 3881 inputs through its real HTTP validation path, from seven generation conditions: two LLM conditions, random values, rule-based adversarial values, boundary values, and two reference conditions. Semantic validity was decided by deterministic rule functions for objectively decidable rules and by an independent judge model for the rest, with a human annotation of 200 inputs. Three results do not depend on the judge. Rules the application could enforce in code were broken by 7.9 % of structurally valid LLM inputs, and by 1.8 % once a target whose violations come from the prompts never stating the current date is removed, which is below the pre-registered bound. Random structurally valid inputs broke rules more often than LLM inputs (53.3 % against 30.9 %), so the gap is a property of the unenforced boundary rather than of AI writers. And a rule layer costs at most 14 ms of p97.5 latency, an embedding layer 65–96 ms, while an LLM judge costs seconds per request and saturates on two of three endpoints at eight concurrent clients. Results that rest on judgement rules are exploratory: calibrated against designed conditions the judge caught 78.6 % of planted violations but flagged 14.1 % of valid inputs, which reduces the judgement-rule rate from 15.7 % to 2.6 % once corrected, and agreement with a human annotator was  $\kappa = 0.24$ . None of five candidate validation layers reached the pre-registered target of 70 % recall at a 5 % false-positive rate. The study covers one application and CPU-only open models.

**Index Terms**—input validation, business logic, large language models, web applications, software testing, empirical study

## I. INTRODUCTION

A web application accepts an input only after it passes validation. In the JavaScript and TypeScript ecosystem that validation is usually a schema. The schema is a declarative object that states the type, presence, length, pattern, numeric range, enumeration and format of every field. It is checked on the client, again on the server, and finally by database constraints. Schemas of this kind are well understood [1], [2] and have been studied empirically at scale [3]. Other rules depend on what the application is for: a promotional code has to apply to the item being bought, a workshop cannot start in the past, and a review is supposed to be about the workshop

it is attached to. Some of these can be written as custom refinements, although in practice most are never written down anywhere. The distinction that matters is therefore between the business rules an application enforces and the many it leaves unenforced.

That limitation is not new. Felmetsger et al. separated input-validation flaws, which admit a general specification, from logic vulnerabilities, which do not [4]. Later work detected such flaws from behavioural traces [5], from client–server validation mismatches [6] and from over-permissive API payloads [7]. What has changed since then is who writes the input. Large language models (LLMs) now fill forms, operate web applications on a user’s behalf, and generate test data [8]–[10].

Recent work shows that LLM output which satisfies a schema is often wrong in meaning. Schema-constrained ordering agents reach full schema validity while semantic success falls as low as a few per cent [11]; hard schema decoding raised validity while executable accuracy fell [12]; across 21 models on its text leaderboard, schema compliance exceeds value accuracy by 15 to 25 percentage points [13]; and in a single-model pilot, schema-first tool interfaces shifted failures from interface misuse to schema-valid semantic misuse [14]. In one enterprise platform the most structurally successful model had the lowest semantic satisfaction [15]. All of this concerns model outputs that a tool, an API or a renderer consumes, and there a wrong value is discarded. At the input boundary of a web application the accepted value is stored and later shown to other users or to an assistant, and two questions remain open there. How often do LLM-written inputs that pass validation break the application’s business rules? And is that rate different from the rate for other producers of structurally valid input? Web testing studies cannot answer either, because they treat acceptance by the application as validity [8], [9]; one of them notes that incomplete e-mail addresses may be submitted successfully because the application lacks the check [8], and a robustness study using status codes as its oracle cannot see failures that return success [16]. Where semantic validity is measured, a base eight-billion-parameter model produces domain-coherent values under a quarter of the time [17], with no non-LLM comparison.

This paper measures that gap at the input boundary of one purpose-built application with an explicit rule reference, and compares LLM inputs against random, rule-based, boundary

and reference inputs on the same outcome. I found no prior study that measures semantic validity at a web application’s input boundary against an enumerated rule reference, compares LLM producers with non-LLM producers on that outcome, or reports what a semantic validation layer costs in the request path on commodity hardware. The study establishes five things.

- **F1. Inputs that the application accepts break rules it does not enforce, and the effect is measurable without any model in the loop.** Rules decidable by code alone were broken by 7.9 % of structurally valid model inputs and 10.3 % of random ones.
- **F2. The gap belongs to the validation boundary, not to AI writers.** Random structurally valid inputs broke rules more often than model inputs (53.3 % against 30.9 %), and the ordering holds in every sensitivity variant. Any producer that satisfies the schema without knowing the rules falls into the gap.
- **F3. A judge model is not yet a usable oracle for this task, and its error is measurable.** Calibrated against planted violations and known-valid inputs, the judge caught 78.6 % of violations while flagging 14.1 % of valid inputs, which moves the judgement-rule estimate from 15.7 % to 2.6 % once corrected. The calibration design itself transfers to any study that uses a model as ground truth.
- **F4. No cheap validation layer closes the gap.** Of five designs evaluated at a bounded false-positive rate on held-out fields, the best reached 27.1 % recall, and most of what it caught were violations a line of code could have enforced.
- **F5. On commodity hardware the cost ordering is decisive.** A rule layer adds at most 14 ms to p97.5 latency, an embedding layer 65–96 ms, and an LLM judge seconds per request, saturating on two of three endpoints at eight concurrent clients. The guardrail literature I reviewed reports no comparable CPU figures [18]–[21].

The testbed, harness, rule reference, raw data and analysis scripts are released, including the defects found during execution and how each was handled.

## II. BACKGROUND AND TERMINOLOGY

**Structural validation** is every check the application performs before business logic: the runtime schema (type, presence, length, pattern, range, enumeration, format and any refinements written into it), service-level checks and database constraints. An input *passes* when the application accepts it. **Semantic validity** is whether a structurally valid value is appropriate for the application’s purpose, covering consistency across fields, domain plausibility, compliance with business rules and safety for downstream consumers such as an LLM component. It is operationalised as satisfying every rule in a frozen, enumerated reference.

An input that passes structural validation and violates at least one rule is **SV-SI** (structurally valid, semantically invalid). The *conditional* SV-SI rate is the share of SV-SI

inputs among the structurally valid inputs whose rules were all decided. It is the primary outcome. A generator that produces more well-formed values would otherwise look worse merely by passing more often.

Terminology differs across communities. The OWASP Input Validation Cheat Sheet separates syntactic from semantic validation [22]; I count a range the developer wrote into the schema as structural, because the application enforces it. Security research speaks of logic or business logic vulnerabilities [4], [5], [23], API testing research of inter-parameter dependencies [24], [25]. In LLM security the adjective usually means adversarial intent [26], whereas my usage concerns whether the data is appropriate, which includes that case.

## III. RELATED WORK

### A. Limits of structural validation, and logic flaws

JSON Schema has a formal semantics and a known complexity [1], [2], and empirical studies describe which constraints practitioners encode [3]. Inter-parameter dependencies such as *requires*, *only-one* and arithmetic relations cannot be stated in an OpenAPI document [24], [25], and schema-derived fuzzers inherit the same boundary [27], and surveys of input-validation vulnerabilities remain organised around injection classes [28]. The rules beyond that boundary are the classical subject of logic-flaw research. Waler inferred likely invariants over session state and database query parameters and model-checked them [4], and related classes are detected by behavioural patterns [5], parameter tampering [6], mass-assignment testing [7] and excessive-data-exposure fuzzing [29]. Recent work applies LLMs to find such flaws in code [23]. In the studies I reviewed the inputs are crafted by an attacker or a tester, and they are not compared across producers.

### B. LLM-generated inputs, and schema-valid but wrong outputs

LLMs are better than random and type-based generators at producing inputs that applications accept. On thirty real forms an LLM-driven approach reaches a form passing rate of 83 %, against 23 % for a static type-based generator and 3 % for random crawling [9], and across 146 forms the best of eleven models submits over 98 % of its inputs [8]. The oracle in both cases is acceptance by the application. Where semantic validity is measured separately, a base model reaches 22.9 % domain-coherent values against 72.4 % after fine-tuning [17]. Related work also mines input constraints from specifications and benchmarks form-filling agents [10]. For model *outputs*, several 2026 studies document the same pattern: OrderBench separates syntactic validity, schema validity and semantic success [11]; hard schema decoding costs accuracy [12]; JSON pass rate stays ahead of value accuracy [13]; schemas shift failures from interface misuse to semantic misuse [14]; and workflows pass validation and then fail to execute [15]. Tool-use benchmarks evaluate emitted calls against a task’s expected answer [30], [31]. In all of them the semantic check is the oracle rather than a layer that is itself evaluated.

### C. Data validation beyond schemas, and inappropriate input as a security problem

Data-management research validates data beyond structural schemas with anomaly detection for machine-learning pipelines [32], declarative constraint checks that act as unit tests for data [33] and learned detection of column semantic types [34]; these work on datasets in batch rather than on one request. Classifying a single field entry from a context embedding and statistics of recent entries has been described in a patent [35], of which my EMB design is a simpler, untrained relative. On the security side, indirect prompt injection established that models do not separate data from instructions [36], and that attacks on deployed applications succeed [37]. Prompt fragments inserted through an ordinary form field can later drive a database write [38], the pattern occurs in the wild [39], and benchmarks extend the coverage [40], [41]. Those inputs are attacker-crafted and mostly free text.

### D. Detection layers and their cost

Input-side detectors include fine-tuned classifiers [18], embedding classifiers [19], [42], small transformers [43] and LLM judges [44]. What constrains deployment is the false-positive rate on benign traffic [45]. Cost evidence is thin and tied to GPUs. A detector presented as deployable gives no timing at all [18], an embedding guardrail puts its embedding floor at 35 ms on a datacentre GPU [19], and a layered framework reports a 61 ms median on an A100 with model inference excluded [20]. A systematisation notes that cost and utility are routinely overlooked [21], the one study that compares regex, classifier and LLM tiers head to head does not state its hardware [46], and LLM judges can themselves be redirected by the content they assess [47]. I found no measurement of a validator for data appropriateness on structured request payloads on CPU.

### E. Position of this work

Three lines of work come closest among the studies I reviewed, and the design of this study responds to a specific difference in each. The output-side studies measure schema-valid but wrong values that a model emits [11]–[13], [15], where the oracle is the task’s expected answer; I measure values that an application accepts and stores, where it raises no error and the value is later read by people or by an assistant. The form and API testing studies count acceptance by the application [8], [9], so an accepted but inappropriate value counts as a success, and the one study that measures domain coherence separately has no non-LLM comparison [17]; I separate the two outcomes and label semantic validity against a frozen rule reference, independently of the application. The guardrail studies evaluate detectors for adversarial intent and, where they report cost at all, do so on datacentre GPUs [19], [20]; I evaluate detectors for ordinary data appropriateness, on structured payloads, on a four-core CPU, and report detection quality and cost together. The comparison across seven producers, including non-LLM ones, is what lets me ask whether

the gap belongs to AI writers or to the validation boundary itself.

## IV. RESEARCH QUESTIONS AND HYPOTHESES

The following were frozen, with the primary metrics and the analysis plan, before any experimental input was generated. Prompt family P5 (injection, Section V) is excluded from all primary rates and analysed separately.

- **RQ1.** Do LLM-generated inputs pass structural validation while violating business rules? *H1:* the conditional SV-SI rate of the LLM conditions has a 95 % Wilson interval whose lower bound exceeds 0.05.
- **RQ2.** How does this vary across field categories? *H2:* rates differ across categories; free text and cross-field highest, simple scalar and structured (enumeration) lowest.
- **RQ3.** How do LLM inputs compare with other producers? *H3a:* LLM with field context exceeds random; *H3b:* LLM with field context exceeds benign-unusual reference inputs; *H3c:* LLM without field context exceeds random.
- **RQ4.** Which violation types dominate? Exploratory.
- **RQ5.** Can an added semantic layer detect SV-SI inputs? *H5:* at least one non-judge design reaches recall  $\geq 0.70$  at a false-positive rate (FPR)  $\leq 0.05$  on held-out fields.
- **RQ6.** What does it cost? *H6:* rule and embedding designs add at most 50 ms to p95 request latency while an LLM judge adds at least 1000 ms.

## V. METHOD

### A. Subject under study

The subject is a purpose-built Next.js 16 application written in TypeScript. It uses Zod 4 for validation on the client and the server, Prisma with SQLite for persistence, and a local Ollama runtime for two AI-facing features, support-ticket triage and a booking assistant. It has seven input surfaces and 39 fields. A reference file enumerates 43 business rules, of which five are enforced in code, two of those as Zod refinements (a birth date not in the future, a workshop end after its start). The remaining 38 are unenforced, a profile chosen to resemble small production applications. Each rule is classified in advance as objective, meaning decidable from the input, the database and the clock (19 rules), or as judgement (24 rules). The application and the rule reference were tagged before any input was generated (commit 62a7267), and the commit is recorded with every run. A purpose-built subject is a deliberate trade: real applications do not enumerate their business rules, and without a complete reference the semantic outcome has no ground truth.

### B. Targets and generation conditions

Seventeen targets covering 22 fields were selected to represent every field category, every directly AI-facing field, and five cross-field pairs whose rule relates the two members. Each target carries a purpose description, its rule identifiers and its structural constraints. Seven conditions produce values for each target: **A** valid reference values; **B** random values

that satisfy the schema, from a seeded generator of word sequences, numbers and dates within the schema’s limits, with no attempt at meaning; **C** rule-based adversarial values, one template per unenforced rule; **D** an LLM given the field name and the structural constraints only; **E** the same LLM given the field’s purpose and rules as well; **F** boundary values at the schema edges; and **G** benign but unusual reference values. Conditions A and G were specified as human-written. In this execution their 272 values were first drafted with the help of a language model, before the main generation, and checked against the schemas and the objective rules. I then went through every value myself, kept 262, replaced 3 with values of my own, deleted 7, and confirmed that the kept values are ones I would have written. The application accepted 0 of the replacements, because all 3 were promotional codes it does not recognise. They are condition A’s only structural rejections, so the reviewed conditions rest on the kept values. I therefore call these conditions human-curated rather than human-written.

### C. LLM generation

Conditions D and E use qwen3:4b (Q4\_K\_M) through Ollama on a four-core CPU, at temperature 0.8, nucleus 0.9, with a fixed seed per run, three runs, and five values per call. Seven prompt families are stored as hashed files: P1 legitimate values, P2 unusual values, P3 values that satisfy the constraints but “would be wrong for the application to accept”, P4 values whose meaning depends on downstream interpretation, P5 injection for AI-facing targets only, P6 boundary values, and P7 Unicode variants. Because P3 asks for violations, results are reported per family as well as pooled. The model identifier, digest, temperature, seed, prompt file and prompt hash are recorded for every value. Failed calls were counted and never regenerated. There were 88 failure records in 52 prompt cells, of which 75 were call errors such as calls aborted at the 120s timeout and 13 were unparseable replies. A further 430 generated values used key names other than the target’s. A harness defect had filled those keys from the form’s default values, so the records measure the defaults rather than the model, and they are treated as generation failures.

### D. Structural measurement

Each input is measured twice. It is checked against the same schema object the application imports, and it is submitted over HTTP to the real route handler, which records the status, the service errors and the time. The structural outcome is the application’s own decision, so an input rejected by one of the five enforced business checks counts as a structural rejection. Booking state is reset to the seed before each booking request.

### E. Semantic ground truth

Objective rules are decided by rule functions. The 11 objective rules that apply to the 17 targets all have one. Judgement rules are decided by a judge model from a different family than the generator (gemma3:4b, temperature 0). The judge sees the field purpose, the rule text and the submitted value of the target’s fields. It never sees the condition, the

prompt family or the intent. The judge does not see the rest of the record. Seven judgement rules depend on other fields, such as capacity being plausible for the location, and those were judged without that context. A sensitivity analysis excludes them. A human annotator labelled 200 structurally valid inputs blind to condition, with 20 items repeated to estimate the annotator’s consistency. Where a human verdict exists it takes precedence over the judge. An input is *decided* when every applicable rule has a verdict. For compute reasons the judge was applied to a stratified sample: 1200 structurally valid inputs drawn round-robin across target and condition with a fixed seed, plus 136 inputs for two booking targets that were re-measured after a defect (Section IX), up to 14 per stratum and as many as were available.

### F. Analysis population and statistics

Of the 3881 inputs, the following are excluded: 430 key-mismatch records, 131 records produced by inadvertent retries of failed calls, 74 injection-family inputs, and 10 reference values that the reviewer deleted or replaced. That leaves 3236 inputs. Of these, 2917 passed structural validation and 1876 are decided. The remaining passes were either not in the judge sample (1001) or received no usable verdict (40). Rates carry Wilson 95 % confidence intervals (CIs), and contrasts are reported as odds ratios (ORs). H1 is evaluated on the lower bound of its interval, H2 with a  $\chi^2$  test and Cramér’s  $V$ , and H3a to H3c with three one-sided Fisher tests under Holm correction. The judge sample gives the conditions different target mixes, so each contrast is also tested within target with the Cochran–Mantel–Haenszel (CMH) test, which was added after data collection. Values generated in one call, and targets within a category, are not independent, so Wilson intervals are too narrow. Every main rate and contrast therefore also carries a cluster bootstrap interval (2000 draws, resampling whole generation calls and, separately, whole targets; 502 calls and 17 targets). Every other analysis is labelled exploratory. It was fixed in advance that if judge–human agreement on judgement rules falls below  $\kappa = 0.4$ , every result that depends on those rules is reported as exploratory. Every number in this paper is generated from the raw data by scripts released with it. The hypotheses, metrics and analysis plan were fixed in a dated decision log before generation began and are released unchanged. Version control over the analysis code was started after data collection, so that log and the frozen subject tag are the evidence of pre-registration, not a commit history.

## VI. SEMANTIC VALIDATION LAYERS

Five designs run as a drop-in check after structural validation. **R** packages the objective rule functions; because those functions also define the ground truth for objective rules, **R** cannot produce a false positive on them and serves as a floor rather than a competitor. **EMB** scores a value by its cosine distance from an embedding (nomic-embed-text) of the field’s purpose and rules, using no labelled examples, unlike the trained classifier the plan anticipated. **SLM** prompts a small language model (SLM; qwen3:1.7b) with the field purpose and

rules for a verdict and a confidence. **HYB** runs R first and then EMB on what R passes. **JUDGE** sends the same prompt to qwen3:4b, the generator’s model, and serves as a reference point, so SLM and JUDGE receive a rubric similar to the ground-truth judge’s. The prompted designs run at temperature 0, and a reply without a verdict counts as a pass, because the deployed layer fails open.

Evaluation is leave-fields-out over five folds. In each fold the threshold is the smallest score at which at most 5 % of the valid training-fold inputs are flagged, and that threshold is applied to the held-out fields. Designs are compared on recall at that operating point, together with FPR, precision and the threshold-free area under the receiver operating characteristic (ROC) curve (AUROC). The model-based designs were scored on a stratified sample of 575 decided inputs, of which 236 are SV-SI, and all designs are compared on that sample. The fast designs were also scored on all 1937 decided inputs, a population that includes the injection family and the retried records.

## VII. RESULTS

### A. Rate of semantic violations (RQ1)

Among decided, structurally valid LLM inputs, 338 of 1094 (30.9 %, 95 % CI 28.2–33.7) broke at least one business rule (Table I), so H1 is supported. The pooled rate includes prompt family P3, which asks for rule violations. Per family (Table II), the rate was 20.3 % for legitimate-value prompts and 43.3 % when violations were requested. Without P3 it was 28.7 %. Agreement between the judge and the human annotator is below the pre-registered threshold (Section VII E), so every rate that depends on judgement rules is exploratory, this one included. The objective-rule rate below depends only on the rule functions and is not affected. The rate survives most sensitivity variants in Table III, including objective rules alone (7.9 %, CI 6.8–9.2), which do not depend on the judge. It does not survive the strictest one. One target dominates the objective violations: on workshop dates (F15+F16) nearly every model value was in the past, because the prompts never state the current date, which is a property of the harness rather than of the application. Removing both the judge and that target leaves 1.8 % (CI 1.3–2.5), whose lower bound is below the 0.05 threshold, so **H1 is not supported on the one variant that depends neither on the judge nor on the date artefact**. H1 is therefore supported as pre-registered and on judgement-rule evidence, and not supported on judge-independent evidence alone. Per run the rate was 31.1, 41.6 and 40.9 % without field context and 25.7, 24.9 and 22.9 % with it; including the retried records gives 30.5 %. The reference conditions bound the labelling from both sides: valid inputs were labelled invalid in 3.8 % of cases and adversarial inputs in 69.1 %, so about a third of the designed violations were missed and a few valid inputs were flagged.

### B. Field categories (RQ2)

Rates differ across categories (Table IV), but H2 is counted as *not supported*: its directional prediction was that simple

scalar and structured fields would be lowest, and structured came out highest. The association statistic ( $\chi^2(6) = 184.5$ , Cramér’s  $V = 0.31$ ) is reported descriptively only, because with one to four targets per category it tests heterogeneity between targets rather than between categories. These rates also rest on judgement rules, so they are exploratory. The pattern is as follows. Free text (51 %) and cross-field (45 %) are high, as expected, structured fields are the highest (57 %) rather than among the lowest, and financial fields are the lowest (12 %). Categories contain between one and four targets, and per-target rates are in the repository. The structured category is a single target, a ticket’s priority and category, whose rules the judge decided without seeing the ticket text, so the reversal is not a property of enumerated fields in general.

### C. Comparison of producers (RQ3)

None of the three pre-registered contrasts is supported (Table V), and two point the other way. They are computed on labels that include judgement rules, so under the agreement gate they are exploratory as well. The objective-rule sensitivity row shows the same direction without the judge. Random structurally valid inputs broke rules in 53.3 % of cases, more often than LLM inputs without field context (37.7 %) or with it (24.5 %), and that holds in every sensitivity variant (Table III). Clustering changes how firm each contrast is. Resampling whole targets, the difference between context-aware model inputs and random inputs still excludes zero (−51.2—6.6), while the difference between context-free model inputs and random inputs does not (−32.8—2.3), and the pooled LLM rate has a target-clustered interval of 14.1–51.6, far wider than its Wilson interval. Within target (CMH, post hoc) the odds of a violation were lower for LLM inputs with context than for random inputs (odds ratio 0.16, 95 % CI 0.11–0.23) and lower without context (0.47, 0.34–0.66); adversarial inputs exceeded random ones (2.18, 1.20–3.95), and field context lowered the odds relative to the schema alone (0.23, 0.15–0.35). Condition E is given the very rules that define the outcome, so that reduction is an upper bound on what context can achieve rather than a deployable intervention: an application able to put its rules in a prompt can also enforce the objective ones in code. LLM inputs with context and benign-unusual reference inputs did not differ detectably (1.32, 0.77–2.27,  $p = 0.29$ ); the interval is wide, and the reference inputs are AI-drafted values curated by a human reviewer, so this compares model input with human-curated input rather than with input a person wrote unaided.

An exploratory variation repeated conditions D and E on ten targets with families P1 to P4, changing only the generator. The 1.7-billion-parameter model produced more violations than the 4-billion model, both on objective rules (4.8 % against 1.1 %, odds ratio 4.4,  $p_{\text{Holm}} = 0.001$ ) and on all rules (35.0 % against 22.3 %,  $p_{\text{Holm}} = 0.02$ ). Neither the 8-billion model (23.1 %) nor temperature 0 (26.3 %) differed detectably. With one run per generator and a smaller judge sample, this only indicates that a smaller model may widen the gap.

TABLE I

STRUCTURAL AND SEMANTIC OUTCOMES BY INPUT-GENERATION CONDITION. THE CONDITIONAL SV-SI RATE IS THE SHARE OF STRUCTURALLY VALID INPUTS THAT VIOLATE AT LEAST ONE BUSINESS RULE; IT IS THE PRIMARY OUTCOME, COMPUTED OVER *decided* PASSES: THOSE WHOSE EVERY APPLICABLE RULE RECEIVED A VERDICT (PASSES OUTSIDE THE JUDGE SAMPLE ARE UNDECIDED AND EXCLUDED). PROMPT FAMILY P5 (INJECTION) AND RETRIED RECORDS ARE EXCLUDED. INTERVALS ARE WILSON 95% CONFIDENCE INTERVALS. \*SPECIFIED AS HUMAN-WRITTEN; VALUES AI-DRAFTED AND THEN REVIEWED VALUE BY VALUE BY THE RESEARCHER (HUMAN-CURATED).

| Condition                | $n$  | Struct. pass (%) | Passes | Decided | SV-SI | Cond. SV-SI (%) | 95% CI       |
|--------------------------|------|------------------|--------|---------|-------|-----------------|--------------|
| A valid human*           | 134  | 97.8             | 131    | 130     | 5     | 3.8             | [1.7, 8.7]   |
| B random                 | 425  | 93.6             | 398    | 362     | 193   | 53.3            | [48.2, 58.4] |
| C rule-based adversarial | 85   | 96.5             | 82     | 81      | 56    | 69.1            | [58.4, 78.1] |
| D LLM context-free       | 1268 | 84.0             | 1065   | 530     | 200   | 37.7            | [33.7, 41.9] |
| E LLM context-aware      | 1108 | 92.8             | 1028   | 564     | 138   | 24.5            | [21.1, 28.2] |
| F boundary               | 85   | 96.5             | 82     | 81      | 26    | 32.1            | [22.9, 42.9] |
| G benign-unusual human*  | 131  | 100.0            | 131    | 128     | 22    | 17.2            | [11.6, 24.7] |

TABLE II

CONDITIONAL SV-SI RATE OF MODEL-GENERATED INPUTS BY PROMPT FAMILY (DECIDED STRUCTURAL PASSES; P5 EXCLUDED). P3 EXPLICITLY ASKS THE MODEL FOR VALUES THE APPLICATION SHOULD REJECT; P1 ASKS FOR LEGITIMATE VALUES. D: WITHOUT FIELD CONTEXT; E: WITH PURPOSE AND RULES.

| Prompt asks for      | D (%) | E (%) | SV-SI    | D+E (%) | 95% CI       |
|----------------------|-------|-------|----------|---------|--------------|
| P1 legitimate        | 25.9  | 14.4  | 43/212   | 20.3    | [15.4, 26.2] |
| P2 unusual           | 36.2  | 26.1  | 52/163   | 31.9    | [25.2, 39.4] |
| P3 rule violations   | 48.1  | 39.1  | 71/164   | 43.3    | [35.9, 50.9] |
| P4 context-dependent | 29.1  | 20.6  | 47/193   | 24.4    | [18.8, 30.9] |
| P6 boundary          | 46.7  | 22.3  | 56/169   | 33.1    | [26.5, 40.5] |
| P7 Unicode variation | 45.6  | 27.2  | 69/193   | 35.8    | [29.3, 42.7] |
| All (H1)             |       |       | 338/1094 | 30.9    | [28.2, 33.7] |
| All except P3        |       |       | 267/930  | 28.7    | [25.9, 31.7] |

TABLE III

SENSITIVITY OF THE CONDITIONAL SV-SI RATE TO DESIGN CHOICES. LLM: CONDITIONS D AND E POOLED; B: RANDOM; G: BENIGN-UNUSUAL (HUMAN-CURATED). "WITHOUT CONTEXT RULES" DROPS SEVEN JUDGEMENT RULES THE JUDGE HAD TO DECIDE WITHOUT THE REST OF THE RECORD. THE TWO OBJECTIVE ROWS USE THE RULE FUNCTIONS ALONE AND SO DO NOT DEPEND ON THE JUDGE. TARGET-WEIGHTED AVERAGES PER-TARGET RATES BY EACH TARGET'S SHARE OF STRUCTURAL PASSES. P1 AND P3 APPLY TO MODEL CONDITIONS ONLY.

| Variant                          | LLM (%) | 95% CI       | B (%) | G (%) |
|----------------------------------|---------|--------------|-------|-------|
| Pre-registered (all rules)       | 30.9    | [28.2, 33.7] | 53.3  | 17.2  |
| P1 (legitimate) prompts only     | 20.3    | [15.4, 26.2] | —     | —     |
| Without P3 (asks for violations) | 28.7    | [25.9, 31.7] | 53.3  | 17.2  |
| Without dates (F15+F16)          | 21.5    | [19.0, 24.2] | 50.4  | 17.5  |
| Objective rules only (no judge)  | 7.9     | [6.8, 9.2]   | 10.3  | 0.8   |
| Objective only, without dates    | 1.8     | [1.3, 2.5]   | 5.1   | 0.0   |
| Without context rules            | 24.5    | [22.1, 27.2] | 36.5  | 9.4   |
| Target-weighted                  | 30.2    | —            | 53.3  | 17.2  |

#### D. Violation types (RQ4)

Counts of judgement-rule violations depend on the judge and are exploratory. Judgement-rule violations outnumber objective ones (502 against 272, across all conditions). The single most frequent violation is objective: a workshop must start in the future (146). It is followed by a profile bio addressed to the assistant (64), a ticket priority inconsistent with its text (61) and duration plausibility (57). The ticket rule was judged without the text it depends on. On AI-facing fields,

TABLE IV

CONDITIONAL SV-SI RATE BY INPUT-FIELD CATEGORY, OVER DECIDED STRUCTURAL PASSES FROM ALL CONDITIONS. A  $\chi^2$  TEST OF INDEPENDENCE GIVES  $\chi^2 = 184.5$ ,  $p < 0.001$ , Cramér's  $V = 0.314$ .

| Field category | Decided | SV-SI | Cond. SV-SI (%) | 95% CI       |
|----------------|---------|-------|-----------------|--------------|
| Structured     | 107     | 61    | 57.0            | [47.5, 66.0] |
| Free text      | 69      | 35    | 50.7            | [39.2, 62.2] |
| Cross-field    | 600     | 270   | 45.0            | [41.1, 49.0] |
| AI-facing      | 260     | 114   | 43.8            | [37.9, 49.9] |
| Search         | 95      | 30    | 31.6            | [23.1, 41.5] |
| Simple scalar  | 249     | 68    | 27.3            | [22.2, 33.2] |
| Financial      | 496     | 62    | 12.5            | [9.9, 15.7]  |

TABLE V

PRE-DECLARED ONE-SIDED CONTRASTS ON THE CONDITIONAL SV-SI RATE (FISHER EXACT, HOLM-CORRECTED ACROSS THE THREE TESTS). EACH TEST IS ONE-SIDED IN THE PRE-REGISTERED DIRECTION (RATE 1 > RATE 2), SO AN ODDS RATIO BELOW 1 YIELDS  $p$  CLOSE TO 1. CONDITIONS: B RANDOM, E LLM WITH FIELD CONTEXT, G BENIGN-UNUSUAL (HUMAN-CURATED), D LLM WITHOUT FIELD CONTEXT. WITHIN-TARGET RESULTS ARE IN THE TEXT.

| Contrast | Rate 1 (%) | Rate 2 (%) | OR   | $p_{\text{Holm}}$ | Supported |
|----------|------------|------------|------|-------------------|-----------|
| H3a E>B  | 24.5       | 53.3       | 0.28 | 1.00              | no        |
| H3b E>G  | 24.5       | 17.2       | 1.56 | 0.14              | no        |
| H3c D>B  | 37.7       | 53.3       | 0.53 | 1.00              | no        |

81 structurally valid LLM inputs violated a rule. Of those, 20 came from the injection family and 61 from the other prompt families.

#### E. How good is the ground truth?

Two checks bound the label quality, and they matter more than any single rate below.

*Calibration against the designed conditions.* Condition C is built from templates that each break one named rule, and condition A is written to be valid, so for those inputs the truth is known without any annotator. On judgement rules the judge caught 44 of 56 designed violations (recall 78.6%, CI 66.2–87.3) and flagged 19 of 135 valid inputs (false-positive rate 14.1%, CI 9.2–20.9). Per-rule figures are in Table VI. A false-positive rate of that size matters, because the observed judgement-rule rate among model inputs is only 15.7%: correcting it for the measured recall and false-positive

TABLE VI

JUDGE CALIBRATION ON THE TWO DESIGNED CONDITIONS, BY JUDGEMENT RULE (RULES WITH AT LEAST TWO DESIGNED VIOLATIONS OR AT LEAST ONE FALSE POSITIVE, ORDERED BY THE NUMBER OF DESIGNED VIOLATIONS). RECALL IS MEASURED ON CONDITION C, WHOSE TEMPLATES EACH BREAK ONE NAMED RULE; THE FALSE-POSITIVE RATE IS MEASURED ON CONDITION A, WHICH IS WRITTEN TO BE VALID. FULL TABLE IN THE REPOSITORY.

| Rule   | Caught | Recall (%) | False pos. | FPR (%) |
|--------|--------|------------|------------|---------|
| B-CO-7 | 5/5    | 100.0      | 7/8        | 87.5    |
| B-BK-7 | 5/5    | 100.0      | 2/8        | 25.0    |
| B-PR-6 | 5/5    | 100.0      | 1/8        | 12.5    |
| B-CO-4 | 4/5    | 80.0       | 0/8        | 0.0     |
| B-PR-5 | 5/5    | 100.0      | 0/7        | 0.0     |
| B-CO-5 | 2/5    | 40.0       | 0/8        | 0.0     |
| B-SE-3 | 5/5    | 100.0      | 0/8        | 0.0     |
| B-TK-3 | 1/3    | 33.3       | 0/8        | 0.0     |
| B-CH-2 | 2/3    | 66.7       | 0/8        | 0.0     |
| B-TK-5 | 1/3    | 33.3       | 0/8        | 0.0     |

rate gives 2.6%. Condition C is deliberately blatant, so its recall is optimistic and the corrected figure is a lower bound, but the gap between 15.7% and 2.6% shows that most of the judgement-rule signal in this study is judge error rather than violation.

*Agreement with a human.* On judgement rules the judge agreed with the annotator on 70.6% of 126 verdicts (Cohen’s  $\kappa = 0.24$ ), and the annotator was unsure on 3.6%. Agreement is below the pre-registered 0.4, so every judgement-rule result in this paper is exploratory. That annotation cannot settle which rater is at fault, because the same annotator also disagreed with the deterministic rule functions on objective rules (89.5% agreement,  $\kappa = 0.32$ , 19 rule-function violations against 4), and those rules are decidable by code. A single annotator working at 3 s per item is therefore not an adequate reference, which is why the calibration above carries the argument instead. The annotator was self-consistent on repeated items ( $\kappa = 0.71$ ).

#### F. Detection quality (RQ5)

No design reached the target, so H5 is not supported (Table VII). Detection is measured against the same labels, so the judgement-rule part of these results is exploratory as well. The objective-rule part is not. At the operating point fixed on the training folds, the best non-judge design, HYB, detected 27.1% of SV-SI inputs (CI 21.3–32.5) at a held-out FPR of 5.6%, which is above the 5% it was tuned for. Most of what it catches comes from R. HYB and R both caught 100.0% of the inputs that violate only objective rules, and 6.5% and 0.0% of those that violate a judgement rule. On all decided inputs the fast designs reach 33.3% for R and 36.6% for HYB. EMB on its own was no better than chance (AUROC 0.45).

The prompted designs rank inputs better (AUROC 0.75 for SLM and 0.76 for JUDGE against 0.60 for HYB), but their scores take only a handful of values, so holding FPR to 5% on the training fields forces the threshold to the top confidence level, where SLM detects 1.3% and JUDGE 23.7%. Taking

each model’s own verdict without tuning, which is exploratory, SLM detects 69.1% while flagging 20.1% of valid inputs, and JUDGE detects 51.7% while flagging 6.8% and gives no usable verdict for 26% of inputs. Both receive a rubric similar to the ground-truth judge’s, so on judgement rules their measured “recall” is largely agreement between small language models and the small model that produced the labels, not detection of violations. Together with R’s tautological recall on objective rules, this means RQ5 evaluates weak baselines: a trained classifier, which the plan anticipated and the leave-fields-out labels would support, was not built, so C3 should be read as a negative result about untrained designs. Thresholds did not transfer across fields; the held-out FPR per fold reached 30% for the embedding-based designs. On benign-unusual inputs the FPR was 4.8% for HYB, 1.6% for SLM and 1.6% for JUDGE.

#### G. Overhead (RQ6)

Each configuration ran against the production build with only its own model loaded. No request failed and no layer call failed in any run (Table VIII). H6 is not supported, because its first part fails. R added at most 1 ms to p97.5 latency with one client and 14 ms with eight, which is within the 50 ms bound. The negative differences at eight clients reflect run-to-run variation. EMB added 65–96 ms with one client, above the bound, and 122–751 ms with eight, because every request embeds its text on the same CPU. On the p95 bracket used as a sensitivity check, EMB fails on the review endpoint and is undetermined on the other two. The second part of H6 holds. SLM added 527–534 ms at the median with one client and JUDGE added 1611–7704 ms. The model server answers one request at a time, so requests queue under concurrency. With eight clients SLM added 4149–4210 ms at the median, and JUDGE completed no request within the 60 s window in all three runs on the course endpoint and in two of three on the review endpoint. Throughput on the course endpoint fell from 47.4 requests per second with structural validation alone to 1.8 with SLM and 0.1 with JUDGE. Set against detection quality (Fig. 1), only R and HYB lie on the trade-off frontier. SLM, JUDGE and EMB each cost more and detect less than one of them at the pre-registered operating point.

## VIII. DISCUSSION

Throughout this section, every statement that rests on judgement rules is exploratory, because judge–human agreement fell below the pre-registered threshold. Statements about objective rules do not depend on the judge.

**RQ1: the gap is measurable, and its size depends on what is asked.** About a third of structurally valid LLM inputs broke an unenforced rule. The figure is a fifth when the model was asked only for legitimate values, and under a tenth on objective rules alone. This extends output-side findings [11], [13], [15] to inputs that an application would store. The magnitude is sensitive to prompt intent, to one date-dependent target, and to the judge, whose agreement with a human annotator was  $\kappa = 0.24$ . Two checks it does not survive

TABLE VII

DETECTION QUALITY OF THE CANDIDATE SEMANTIC VALIDATION LAYERS ON THE SHARED STRATIFIED SAMPLE OF 575 DECIDED INPUTS (236 SV-SI, 339 VALID), LEAVE-FIELDS-OUT (5 FOLDS). RECALL, 95% BOOTSTRAP INTERVAL, FPR AND PRECISION ARE AT THE PRE-REGISTERED OPERATING POINT: THE THRESHOLD IS THE SMALLEST SCORE WHOSE FALSE-POSITIVE RATE ON VALID TRAINING-FOLD INPUTS IS AT MOST 5%, APPLIED TO HELD-OUT FIELDS. AUROC IS THRESHOLD-FREE. THE NATIVE COLUMN (EXPLORATORY, NOT PRE-REGISTERED) APPLIES EACH DESIGN'S OWN DECISION RULE WITHOUT TUNING: FOR SLM AND JUDGE, THE MODEL ANSWERED "NOT APPROPRIATE"; FOR R, EMB AND HYB, SCORE  $\geq 0.5$ . JUDGE RETURNED NO USABLE VERDICT FOR 26% OF INPUTS, COUNTED AS NOT FLAGGED. R'S ZERO FPR IS PARTLY BY CONSTRUCTION: ITS RULE FUNCTIONS ALSO DECIDE THE GROUND TRUTH FOR OBJECTIVE RULES. JUDGE SHARES A MODEL FAMILY WITH THE GENERATOR AND IS A REFERENCE POINT. LATENCY IS THE MEDIAN DECISION TIME OF THE LAYER ALONE ON THE STUDY CPU; HYB'S FIGURE REUSES EMBEDDINGS ALREADY CACHED BY EMB, SO ITS UNCACHED REQUEST-PATH COST IS THE ONE IN TABLE VIII.

| Design | Recall (%) | 95% CI       | FPR (%) | Precision (%) | AUROC | Native recall / FPR (%) | Median (ms) |
|--------|------------|--------------|---------|---------------|-------|-------------------------|-------------|
| HYB    | 27.1       | [21.3, 32.5] | 5.6     | 77.1          | 0.60  | 24.2 / 2.1              | 0.01        |
| JUDGE  | 23.7       | [18.5, 29.4] | 3.5     | 82.4          | 0.76  | 51.7 / 6.8              | 2709        |
| R      | 22.5       | [17.0, 27.8] | 0.0     | 100.0         | 0.61  | 22.5 / 0.0              | 0.01        |
| EMB    | 4.7        | [2.2, 7.3]   | 5.6     | 36.7          | 0.45  | 1.7 / 2.1               | 31          |
| SLM    | 1.3        | [0.0, 2.9]   | 1.2     | 42.9          | 0.75  | 69.1 / 20.1             | 1075        |

TABLE VIII

REQUEST-PATH OVERHEAD OF EACH CONFIGURATION (PRODUCTION BUILD, FOUR-CORE CPU, NO GPU). EACH VALUE IS THE MEDIAN OVER THREE 60 S RUNS, GIVEN AS THE RANGE ACROSS THE THREE ENDPOINTS (PROFILE, REVIEW, COURSE). ADDED LATENCY IS RELATIVE TO STRUCTURAL-ONLY VALIDATION ON THE SAME ENDPOINT AND CONCURRENCY. H6 IS EVALUATED ON THE ADDED P97.5, THE NEAREST TAIL PERCENTILE THE LOAD GENERATOR REPORTS. "OK" IS THE SHARE OF LAYER CALLS THAT RETURNED A VERDICT RATHER THAN FAILING OPEN. ">60 000": THE CONFIGURATION SATURATED, COMPLETING NO REQUEST WITHIN AT LEAST TWO OF THE THREE RUNS. PER-ENDPOINT VALUES ARE IN THE REPOSITORY.

| Config.              | $\Delta p50$ (ms) | $\Delta p97.5$ (ms) | req/s     | Ok (%) | H6    |
|----------------------|-------------------|---------------------|-----------|--------|-------|
| <i>Concurrency 1</i> |                   |                     |           |        |       |
| structural           | 0                 | 0                   | 42.9–47.4 | –      | –     |
| R                    | 0–2               | 0–1                 | 41.9–46.0 | 100    | holds |
| EMB                  | 27–48             | 65–96               | 13.2–19.2 | 100    | fails |
| HYB                  | 26–45             | 65–82               | 13.6–19.2 | 100    | –     |
| SLM                  | 527–534           | 655–761             | 1.7–1.8   | 100    | –     |
| JUDGE                | 1611–7704         | 7683–10824          | 0.1–0.5   | 100    | holds |
| <i>Concurrency 8</i> |                   |                     |           |        |       |
| structural           | 0                 | 0                   | 41.7–48.3 | –      | –     |
| R                    | 2–6               | -136–14             | 43.9–47.6 | 100    | holds |
| EMB                  | 126–288           | 122–751             | 16.0–27.9 | 100    | fails |
| HYB                  | 103–270           | 131–396             | 17.7–29.0 | 100    | –     |
| SLM                  | 4149–4210         | 6780–7180           | 1.7–1.7   | 100    | –     |
| JUDGE                | 14282–>60 000     | 23474–>60 000       | 0.0–0.4   | 100    | holds |

matter as much as the headline. With the judge removed and the date target removed the rate is 1.8%, below the pre-registered bound. Correcting the judgement-rule component for the judge's measured false-positive rate brings it from 15.7% to 2.6%. The honest summary is that unenforced rules are broken often enough to matter, that this study does not establish the size of the effect, and that the ground truth is the limiting factor.

**RQ2: where it occurs is uneven.** Category is moderately associated with the outcome, but individual targets drive that association. One target with a clock-dependent rule accounts for over half of the cross-field violations, and the structured category is a single target that the judge scored without its context.

**RQ3: the gap is not specific to AI.** This is the central negative result. Random values that satisfy the schema broke




Fig. 1. Detection versus request-path cost for each semantic layer: recall at  $FPR \leq 0.05$  on held-out fields (Table VII) against the largest added p97.5 latency over the three endpoints (Table VIII), per concurrency level. Filled markers lie on the trade-off frontier. The dashed line is the pre-registered 70% recall target. JUDGE at concurrency 8 saturated and is drawn at its lower bound.

rules more often than model values, with or without field context, both pooled and within target. The results point to the validation boundary rather than to the producer. A schema admits many values the application should reject, and any producer that satisfies the schema without knowing the rules will supply some of them. What distinguishes the model is that context helps. Giving it the purpose and the rules cut the odds of a violation by 77% relative to the schema alone. The random generator produces word sequences with no attempt at meaning, so it is a weak producer chosen to bound the space rather than to represent a person, and H3a and H3c were therefore unlikely predictions whose failure carries little information by itself. What the comparison does establish is that satisfying a schema without knowing the rules suffices to produce violations, which is a property of the boundary. The comparison that would settle the AI question, model input against input written unaided by several people, has not been run.

**RQ4: most violations are ordinary.** The single most frequent violation was of an objective rule that the application could enforce directly, so enforcing objective rules would remove part of the gap. Judgement-rule violations are still the

majority. On AI-facing fields most LLM violations came from ordinary prompts rather than from injection attempts, so the gap is mostly unenforced business logic rather than attack.

**RQ5: no cheap layer closes it.** The fast designs that can be held to a low FPR, R and HYB, find about a quarter of the violations at the operating point, and most of what they find are objective violations that could simply be enforced. JUDGE at a similar FPR also finds about a quarter, split evenly between objective and judgement rules. The prompted models discriminate better but produce coarse, field-dependent scores. Used as they are, they catch more while rejecting between one valid input in fifteen and one in five in this sample. Detecting judgement-rule violations remains the open problem.

**RQ6: cost rises faster than detection.** On commodity CPU hardware only the rule layer is effectively free, and it detects violations of rules the application could have enforced in the first place. The embedding layer costs tens of milliseconds per request and detected almost nothing here. The prompted models cost half a second to several seconds per request with one client, a local model server handles one request at a time, so they cannot keep up with eight, and their better ranking of inputs does not survive the false-positive constraint. Two caveats bound that conclusion. Run-to-run variation at eight clients reaches the same order as the embedding effect that fails H6 at one client, so the embedding verdict is marginal. The serialisation is also a property of one local model server: a hosted judge would cost roughly a second per call and scale horizontally, so “too expensive” describes local CPU inference rather than LLM validation in general. On this hardware the practical recommendation is to enforce objective rules in code and to treat judgement rules as an open problem, rather than to add a model to the request path.

**Alternative explanations.** The six explanations fixed in advance fare as follows. *Weak business rules rather than AI:* consistent with the data, because the gap appears for every producer that does not know the rules. *Random inputs produce the same effect:* consistent, and random inputs exceed model inputs. *Unusual human inputs behave the same:* not resolved. Human-curated unusual inputs did not differ detectably from context-aware model inputs (1.32, 0.77–2.27), but the interval spans both directions and the contrast is exploratory. *Confined to free text:* not supported, because structured and cross-field targets also show high rates, although single targets drive them. *Prompt injection under another name:* not supported. Injection prompts are excluded from all rates and account for 20 of 81 LLM violations on AI-facing fields. *Too expensive to deploy:* supported for the model-based layers on CPU hardware, and not supported for rules.

## IX. THREATS TO VALIDITY

**Internal validity.** Labelling is the main threat. Judgement rules were decided by one 4B judge (agreement with a human annotator  $\kappa = 0.24$  on 126 shared verdicts, below the pre-registered 0.4, so judgement-rule results are exploratory), and seven of them were decided without the fields they depend on. Excluding those seven lowers the LLM rate to 24.5 % and

lowers the random rate by more (Table III), which leaves the RQ3 direction intact. Agreement on semantic judgements can be low even between humans [48]. The disagreement here is not one-sided noise. On the shared verdicts the judge called 37 violations against the annotator’s 28, so the judge is the stricter rater and judgement-rule rates may be too high. The judge also returned no usable verdict on 36 further rules that the annotator decided, and the agreement statistic cannot use those. The annotator in turn disagreed with the deterministic rule functions on objective rules ( $\kappa = 0.32$ , 89.5 % agreement, 19 rule-function violations against 4), and those functions decide such rules from the record and the clock. Part of the low agreement is therefore annotation error, or the limited context the annotator was shown, rather than judge error alone. The annotator was self-consistent on the repeated items ( $\kappa = 0.71$ ). The judge sample gives the conditions different target mixes, and target-weighted rates are close to the pooled ones. Values generated in one call are correlated, so the intervals understate uncertainty. Time-relative rules are evaluated against the clock at the moment the labels are merged, so a later re-run can change labels near the boundary. The judge ran once, and its run-to-run stability on this CPU was not measured.

**Defects found and corrected.** Six defects were found during execution, each recorded as a deviation with the invalid data preserved: the database was not reset between booking measurements, so capacity ran out and those targets were re-measured; a resume routine retried failed generation calls, whose records are excluded; model outputs with other key names were silently replaced by form defaults, and those records are excluded as generation failures; a first overhead run never enabled the layer and is discarded; the first detector analysis chose thresholds with a quantile, which fails for coarse scores, and was replaced before any detector result was reported; and the label merge never passed human verdicts to the ground-truth algorithm while the importer dropped cells after an empty one, both fixed before the annotation was collected. Three of the six were found not by my own checking but by separate adversarial review passes over the code and the analysis outputs, run with LLM reviewers.

**External validity.** The study covers one application in one domain, built by the researcher, with one enforcement profile, so the results may not transfer. Generation uses quantised open models of 1.7 to 8 billion parameters on CPU. Frontier models may behave differently, and evidence suggesting that scale does not predict semantic quality [13], [15] cuts both ways. The prompts never state the current date, which inflates date-rule violations for model inputs.

**Construct validity.** Semantic invalidity is defined by a frozen reference that another team would write differently. The reference was frozen before any input existed, and violations are reported per rule, so a reader can discount rules they disagree with.

**Conclusion validity.** Pooling three runs inflates the LLM sample, and per-run rates are reported for that reason. The category mapping assigns dual-tagged fields to one category, which was decided before the data existed but not logged at

the time. The three contrasts H3a to H3c were the only pre-declared ones, but they are computed on labels that include judgement rules, so they too are exploratory under the pre-registered agreement gate, and only results decided by the rule functions alone remain confirmatory.

**Who produced the human data.** I carried out both the review and the annotation, and I also built the subject application and wrote the rule reference. The annotation is blind to the generation condition, because the interface never shows it, but it is not independent of the study, and a single annotator cannot establish how far the disagreement with the judge generalises. The review took 19 minutes for 272 values, a median of 1 s per value, and the annotation took 33 minutes for 200 items, a median of 3 s per item. Both are well below the effort budgeted for them, which may itself contribute to the low agreement. The per-item answers and their timestamps are released with the artefacts.

**Provenance.** The reference conditions A and G were specified as human-written and their first drafts came from a language model. I then reviewed and curated every value, so H3b compares model input with human-curated input rather than with input written unaided. H3b remains exploratory because it depends on judgement rules.

## X. CONCLUSION AND FUTURE WORK

Schema validation answers whether an input is well formed, not whether it belongs. At the input boundary of one application with an explicit rule reference, structurally valid inputs broke unenforced business rules often enough to matter, and random inputs broke them more often than model inputs, so the gap belongs to the boundary rather than to AI writers. How large the gap is remains open: the judge-independent estimate falls below the pre-registered bound once a prompt artefact is removed, and the judgement-rule estimate shrinks by a factor of six once the judge’s false-positive rate is taken into account. The single most frequent violation was of an objective rule that a line of validation code could enforce. None of the five candidate layers detected the pre-registered share of violations at a tolerable false-positive rate. A human annotation of 200 inputs agreed with the judge only at  $\kappa = 0.24$ , below the threshold fixed in advance, so every result that depends on judgement rules is exploratory, and the ground truth for those rules is the weakest part of the study. On a four-core CPU, rules added at most 14 ms at p97.5 and embeddings tens of milliseconds. At the median a small language model added about half a second and an LLM judge several seconds, and the judge saturated on two of three endpoints with eight concurrent clients.

Future work should strengthen the ground truth with several annotators and more time per item, collect reference inputs written unaided by several people, re-judge the context-dependent rules with the full record, and repeat the protocol on an existing open-source application.

## ACKNOWLEDGMENT

The testbed, the harness, the rule reference, every generated input, measurement and label, and the scripts that produce every table, figure and number in this paper are available at <https://github.com/zahidrahimoon/semantic-validation-gap>. Generative AI tools assisted with drafting and code; all content was verified by the author, who is responsible for it.

## REFERENCES

- [1] F. Pezoa, J. L. Reutter, F. Suarez, M. Ugarte, and D. Vrgoč, “Foundations of JSON schema,” in *Proceedings of the 25th International Conference on World Wide Web*, ser. WWW ’16. International World Wide Web Conferences Steering Committee, Apr. 2016, pp. 263–273.
- [2] L. Atouche, M.-A. Baazizi, D. Colazzo, G. Ghelli, C. Sartiani, and S. Scherzinger, “Validation of modern JSON schema: Formalization and complexity,” *Proceedings of the ACM on Programming Languages*, vol. 8, no. POPL, pp. 1451–1481, 2024.
- [3] M.-A. Baazizi, D. Colazzo, G. Ghelli, C. Sartiani, and S. Scherzinger, *An Empirical Study on the “Usage of Not” in Real-World JSON Schema Documents*. Springer International Publishing, 2021, pp. 102–112.
- [4] V. Felmetzger, L. Cavedon, C. Kruegel, and G. Vigna, “Toward automated detection of logic vulnerabilities in web applications,” in *Proceedings of the 19th USENIX Security Symposium*. USENIX Association, 2010. [Online]. Available: [https://www.usenix.org/legacy/event/sec10/tech/full\\_papers/Felmetzger.pdf](https://www.usenix.org/legacy/event/sec10/tech/full_papers/Felmetzger.pdf)
- [5] G. Pellegrino and D. Balzarotti, “Toward black-box detection of logic flaws in web applications,” in *Proceedings 2014 Network and Distributed System Security Symposium*, ser. NDSS 2014. Internet Society, 2014.
- [6] P. Bisht, T. Hinrichs, N. Skrupsky, R. Bobrowicz, and V. N. Venkatakrishnan, “NoTammer: automatic blackbox detection of parameter tampering opportunities in web applications,” in *Proceedings of the 17th ACM conference on Computer and communications security*, ser. CCS ’10. ACM, Oct. 2010, pp. 607–618.
- [7] D. Corradini, M. Pasqua, and M. Ceccato, “Automated black-box testing of mass assignment vulnerabilities in RESTful APIs,” in *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*. IEEE, May 2023, pp. 2553–2564.
- [8] T. Li, C. Cui, R. Huang, D. Towey, and L. Ma, “Large language models for automated web-form-test generation: An empirical study,” *ACM Transactions on Software Engineering and Methodology*, vol. 35, no. 3, pp. 1–37, Feb. 2026.
- [9] P. Alian, N. Nashid, M. Shahbandeh, and A. Mesbah, “Semantic constraint inference for web form test generation,” in *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis*, ser. ISSTA ’24. ACM, Sep. 2024, pp. 932–944, arXiv:2402.00950.
- [10] B. Li, Y. Wang, H. Fei, J. Li, W. Ji, M.-L. Lee, and W. Hsu, “Form-Factory: An interactive benchmarking suite for multimodal form-filling agents,” in *Proceedings of the 33rd ACM International Conference on Multimedia*. ACM, 2025, pp. 13 273–13 280.
- [11] Y. Li, “When JSON is not enough: Semantic reliability of schema-constrained LLM ordering agents,” arXiv preprint arXiv:2607.18261, 2026. [Online]. Available: <https://arxiv.org/abs/2607.18261>
- [12] J. Ray, “The constraint tax: Measuring validity-correctness tradeoffs in structured outputs for small language models,” arXiv preprint arXiv:2605.26128, 2026. [Online]. Available: <https://arxiv.org/abs/2605.26128>
- [13] A. K. Singh, H. V. Khurdula, Y. D. Khemlani, and V. Agarwal, “The structured output benchmark: A multi-source benchmark for evaluating structured output quality in large language models,” arXiv preprint arXiv:2604.25359, 2026. [Online]. Available: <https://arxiv.org/abs/2604.25359>
- [14] A. Sigdel and R. Baral, “Schema first tool APIs for LLM agents: A controlled study of tool misuse, recovery, and budgeted performance,” arXiv preprint arXiv:2603.13404, 2026. [Online]. Available: <https://arxiv.org/abs/2603.13404>

- [15] X. Wrenn, R. Raykov, A. Angelov, H. Kitahara, Y. Watanabe, and A. Sailer, "Evaluating LLM trade-offs for enterprise automation: Lessons from workflow generation in a production enterprise platform," arXiv preprint arXiv:2608.03311, 2026, accepted at the 14th IEEE International Conference on Cloud Engineering (IC2E 2026), Industry Track. [Online]. Available: <https://arxiv.org/abs/2608.03311>
- [16] H. G. Tigulla and M. Vieira, "LLM-based robustness testing of microservice applications: An empirical study," arXiv preprint arXiv:2605.14202, 2026. [Online]. Available: <https://arxiv.org/abs/2605.14202>
- [17] M. Kim, S. Sinha, and A. Orso, "LlamaRestTest: Effective REST API testing with small language models," *Proceedings of the ACM on Software Engineering*, vol. 2, no. FSE, pp. 465–488, Jun. 2025.
- [18] D. Jacob, H. Alzahrani, Z. Hu, B. Alomair, and D. Wagner, "Prompt-Shield: Deployable detection for prompt injection attacks," in *Proceedings of the Fifteenth ACM Conference on Data and Application Security and Privacy*, ser. CODASPY '25. ACM, Jun. 2025, pp. 341–352, arXiv:2501.15145.
- [19] I. Ahmed, A. A. Borsha, R. D. Prangon, A. fuad Ahmad, and T. H. Tran, "Reflex-guard: A low-latency guardrail for LLM prompt safety using dense semantic embeddings," arXiv preprint arXiv:2608.17556, 2026. [Online]. Available: <https://arxiv.org/abs/2608.17556>
- [20] G. Saleem, N. Ahmed, M. I. Zaman, A. Hassan, and U. Mujahid, "A layered security framework against prompt injection in RAG-based chatbots," arXiv preprint arXiv:2606.19660, 2026. [Online]. Available: <https://arxiv.org/abs/2606.19660>
- [21] X. Wang, Z. Ji, W. Wang, Z. Li, D. Wu, and S. Wang, "SoK: Evaluating jailbreak guardrails for large language models," in *2026 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2026, pp. 39–58.
- [22] OWASP Cheat Sheet Series Team, "Input validation cheat sheet," OWASP Cheat Sheet Series, OWASP Foundation, 2026, accessed: Sep. 22, 2026. [Online]. Available: [https://cheatsheetseries.owasp.org/cheatsheets/Input\\_Validation\\_Cheat\\_Sheet.html](https://cheatsheetseries.owasp.org/cheatsheets/Input_Validation_Cheat_Sheet.html)
- [23] B. Metin, M. Wynn, A. Tunali, and Y. Kepir, "Business logic vulnerabilities in the digital era: A detection framework using artificial intelligence," *Information*, vol. 16, no. 7, p. 585, Jul. 2025.
- [24] A. Martin-Lopez, S. Segura, C. Muller, and A. Ruiz-Cortes, "Specification and automated analysis of inter-parameter dependencies in web APIs," *IEEE Transactions on Services Computing*, vol. 15, no. 4, pp. 2342–2355, Jul. 2022.
- [25] A. Martin-Lopez, S. Segura, and A. Ruiz-Cortés, *A Catalogue of Inter-parameter Dependencies in RESTful Web APIs*. Springer International Publishing, 2019, pp. 399–414.
- [26] N. Singh, A. Abuadbbba, Y. Gao, S. Nepal, and H. Kim, "Shifting from injection to interaction: Rethinking web security in the age of LLMs and beyond," arXiv preprint arXiv:2609.03999, 2026. [Online]. Available: <https://arxiv.org/abs/2609.03999>
- [27] Z. Hatfield-Dodds and D. Dygalo, "Deriving semantics-aware fuzzers from web API schemas," arXiv preprint arXiv:2112.10328, 2021, a two-page tool demonstration appeared at ICSE '22 Companion, doi 10.1145/3510454.3528637; this nine-page version is the one cited. [Online]. Available: <https://arxiv.org/abs/2112.10328>
- [28] F. Faisal Fadlalla and H. T. Elshoush, "Input validation vulnerabilities in web applications: Systematic review, classification, and analysis of the current state-of-the-art," *IEEE Access*, vol. 11, pp. 40 128–40 161, 2023.
- [29] L. Pan, S. Cohnney, T. Murray, and V.-T. Pham, "EDEFuzz: A web API fuzzer for excessive data exposures," in *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*, ser. ICSE '24. ACM, Feb. 2024, pp. 1–12.
- [30] P. Yu, W. Liu, Y. Yang, J. Li, Z. Zhang, X. Feng, and F. Zhang, "Benchmarking LLM tool-use in the wild," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2026, arXiv:2604.06185. [Online]. Available: <https://openreview.net/forum?id=yz7fL5vfpn>
- [31] G. Yu, X. Sun, Z. Huang, S. Fan, Z. Zhang, X. Hu, X. Xue, Y. Shi, X. Xiao, Y. Song, B. Lyu, R. Wen, X. Li, Q. He, S. Zhu, and Z. Liu, "Getting the parameters right: A difficulty-graded benchmark and probe-guided training for LLM tool calls," arXiv preprint arXiv:2608.03071, 2026. [Online]. Available: <https://arxiv.org/abs/2608.03071>
- [32] N. Polyzotis, M. Zinkevich, S. Roy, E. Breck, and S. E. Whang, "Data validation for machine learning," in *Proc. Machine Learning and Systems (MLSys)*, vol. 1, 2019. [Online]. Available: [https://proceedings.mlsys.org/paper\\_files/paper/2019/hash/928f1160e52192e3e0017fb63ab65391-Abstract.html](https://proceedings.mlsys.org/paper_files/paper/2019/hash/928f1160e52192e3e0017fb63ab65391-Abstract.html)
- [33] S. Schelter, D. Lange, P. Schmidt, M. Celikel, F. Biessmann, and A. Grafberger, "Automating large-scale data quality verification," *Proc. VLDB Endowment*, vol. 11, no. 12, pp. 1781–1794, 2018.
- [34] M. Hulsebos, K. Hu, M. Bakker, E. Zraggen, A. Satyanarayan, T. Kraska, Ç. Demiralp, and C. Hidalgo, "Sherlock: A deep learning approach to semantic data type detection," in *Proc. 25th ACM SIGKDD Int. Conf. Knowledge Discovery & Data Mining*, 2019, pp. 1500–1508.
- [35] A. Banerjee and V. T. Subrahmaniam, "Deep learning based context embedding approach for detecting data entry errors," US Patent 12 288 159, Apr. 29, 2025. [Online]. Available: <https://patents.google.com/patent/US12288159B2>
- [36] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," in *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, ser. CCS '23. ACM, Nov. 2023, pp. 79–90, arXiv:2302.12173.
- [37] Y. Kaya, A. Landerer, S. Pletinckx, M. Zimmermann, C. Kruegel, and G. Vigna, "When AI meets the web: Prompt injection risks in third-party AI chatbot plugins," in *2026 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2026, pp. 4223–4242.
- [38] R. Pedro, M. E. Coimbra, D. Castro, P. Carreira, and N. Santos, "Prompt-to-SQL injections in LLM-integrated web applications: Risks and defenses," in *2025 IEEE/ACM 47th International Conference on Software Engineering (ICSE)*. IEEE, Apr. 2025, pp. 1768–1780, arXiv:2308.01990.
- [39] S. Khodayari, X. Zhang, B. Acharya, and G. Pellegrino, "Indirect prompt injection in the wild: An empirical study of prevalence, techniques, and objectives," arXiv preprint arXiv:2604.27202, 2026. [Online]. Available: <https://arxiv.org/abs/2604.27202>
- [40] J. Yi, Y. Xie, B. Zhu, E. Kiciman, G. Sun, X. Xie, and F. Wu, "Benchmarking and defending against indirect prompt injection attacks on large language models," in *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.I*, ser. KDD '25. ACM, Jul. 2025, pp. 1809–1820.
- [41] I. Evtimov, A. Zharmagambetov, A. Grattafiori, C. Guo, and K. Chaudhuri, "WASP: Benchmarking web agent security against prompt injection attacks," in *Advances in Neural Information Processing Systems 38*. Neural Information Processing Systems Foundation, Inc., 2025, pp. 22 436–22 458.
- [42] M. A. Ayub and S. Majumdar, "Embedding-based classifiers can detect prompt injection attacks," in *Proceedings of the Conference on Applied Machine Learning in Information Security (CAMLIS 2024)*, ser. CEUR Workshop Proceedings, R. Allen, S. Samtani, E. Raff, and E. Rudd, Eds., vol. 3920. Arlington, Virginia, USA: CEUR-WS.org, 2024, pp. 257–268. [Online]. Available: <https://ceur-ws.org/Vol-3920/paper15.pdf>
- [43] Y. Datta and S. Rajasekar, "JavelinGuard: Low-cost transformer architectures for LLM security," arXiv preprint arXiv:2506.07330, 2025. [Online]. Available: <https://arxiv.org/abs/2506.07330>
- [44] T. Rebedea, R. Dinu, M. N. Sreedhar, C. Parisien, and J. Cohen, "NeMo guardrails: A toolkit for controllable and safe LLM applications with programmable rails," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, Y. Feng and E. Lefever, Eds. Singapore: Association for Computational Linguistics, Dec. 2023, pp. 431–445.
- [45] H. Li and X. Liu, "InjecGuard: Benchmarking and mitigating over-defense in prompt injection guardrail models," arXiv preprint arXiv:2410.22770, 2024. [Online]. Available: <https://arxiv.org/abs/2410.22770>
- [46] A. C. Maiorano, "Evaluating prompt injection defenses for educational LLM tutors: Security-usability-latency trade-offs," arXiv preprint arXiv:2605.06669, 2026. [Online]. Available: <https://arxiv.org/abs/2605.06669>
- [47] J. Shi, Z. Yuan, Y. Liu, Y. Huang, P. Zhou, L. Sun, and N. Z. Gong, "Optimization-based prompt injection attack to LLM-as-a-judge," in *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '24. ACM, 2024, pp. 660–674.
- [48] T. Calò, A.-E. Gurita, and L. De Russis, "Measuring the semantic accessibility gap in LLM-generated web UIs," in *Proceedings of the Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems*, ser. CHI EA '26. ACM, Apr. 2026, pp. 1–5.